

STA 561 Homework 5

April 5, 2023

STA 561 HW 5

Authors

- Alonso Guerrero Castaneda (UID: 1194613)
- Eli Gnesin (UID: 1172961)
- Tommy Misikoff (UID: 1166813)
- Sanskriti Purohit (UID: 1179957)
- Will Tirone (UID: 1130904)

TA: Rick Presman

```
[ ]: # Import Libraries
import numpy as np
import pandas as pd
import random
import numpy.random as nr
from sklearn import datasets
import matplotlib.pyplot as plt
```

Question 1

$$\begin{aligned}\mathbb{E}||k^{\frac{-1}{2}}\Omega\mathbf{x}||^2 &= k^{-1}\mathbb{E}\left\|\begin{bmatrix}\Omega_1^T\mathbf{x} \\ \vdots \\ \Omega_k^T\mathbf{x}\end{bmatrix}\right\|^2 \\ &= k^{-1}\mathbb{E}\left[\left(\Omega_{1,1}x_1 + \dots + \Omega_{1,p}x_p\right)^2 + \dots + \left(\Omega_{k,1}x_1 + \dots + \Omega_{k,p}x_p\right)^2\right] \\ &= k^{-1}\mathbb{E}\left[x_1^2\Omega_{1,1}^2 + \Omega_{1,1}\Omega_{1,2}x_1x_2 + \dots\right] + \dots + \mathbb{E}\left[x_1^2\Omega_{k,1}^2 + \dots\right]\end{aligned}$$

Now the non-squared terms will be 0, since $\mathbb{E}(\Omega_{i,j}) = 0$ and they are i.i.d. For example, $\mathbb{E}(\Omega_{1,1}\Omega_{1,2}x_1x_2) = x_1x_2\mathbb{E}(\Omega_{1,1})\mathbb{E}(\Omega_{1,2}) = 0$

$$\begin{aligned}
&= k^{-1} \left[[x_1^2 \mathbb{E}(\Omega_{1,1}^2) + \dots + x_p^2 \mathbb{E}(\Omega_{1,p}^2)] + \dots + [x_1^2 \mathbb{E}(\Omega_{k,1}^2) + \dots + x_p^2 \mathbb{E}(\Omega_{k,p}^2)] \right] \\
&= k^{-1} \left[[x_1^2 + \dots + x_p^2] + \dots + [x_1^2 + \dots + x_p^2] \right] \quad \text{second moment: } \text{Var}(\Omega_{i,j}) = E(\Omega_{i,j}^2) - 0 = 1 \\
&= k^{-1} \left[\|x\|^2 + \dots + \|x\|^2 \right] \\
&= \|x\|^2
\end{aligned}$$

If we want to generate $\Omega_{i,j} \stackrel{IID}{\sim} Q$, the necessary conditions for the above proof to hold are that Q has a mean of 0 and a variance of 1.

Question 2

$$\begin{aligned}
\hat{\beta}_n^\Omega &= \Omega^T \text{argmin}_\beta \mathbb{P}_n(Y - (\Omega \mathbf{X})^T)^2 \\
&= \Omega^T ((\mathbb{P}_n(\Omega \mathbf{X})(\mathbf{X})^T)^{-1} \mathbb{P}_n(\mathbf{X}) \mathbf{Y}) \\
&= \Omega^T (\Omega \mathbb{P}_n \mathbf{X} \mathbf{X}^T)^{-1} \mathbb{P}_n \mathbf{X} \mathbf{Y}
\end{aligned}$$

Now set $\Sigma = \mathbb{P}_n \mathbf{X} \mathbf{X}^T$ and $\Gamma = \mathbb{P}_n \mathbf{X} \mathbf{Y}$. If our data is very large, say 10 million rows, we can make a single pass through the data to compute these quantities, with the below pseudo code:

```

// initialize
X = data
Y = target
B = number of replicates, chosen by the user
n = length(target)
Sigma = 0
Gamma = 0
Beta = 0

// Loop through and sum
// The goal here is to calculate P_n XX^T and P_n XY exactly once and store them
For i in 1 : n

    // Calculate XX^T and XY
    temp_data = X[i]
    temp_target = Y[i]
    temp_Sigma = temp_data * temp_data^T
    temp_Gamma = temp_data * temp_target

    Sigma = Sigma + temp_Sigma
    Gamma = Gamma + temp_Gamma

Sigma = (1/n) * Sigma
Gamma = (1/n) * Gamma

```

```

// Now loop through our B and calculate beta_hat
For b in 1 : B
    Omega = (k * p matrix randomly sampled from Q)
    Inter = Omega @ Sigma @ Omega^T
    Inv_inter = Inter^{-1}
    projected_beta = Omega^T @ Inv_inter @ Omega @ Gamma

    Beta = Beta + projected_beta

Beta = (1/B) * Beta

// RETURN Beta

```

Regardless of the size of our dataset, we calculate Σ and Γ by iterating through the data once, taking $\mathbf{X}\mathbf{X}^T \in \mathbb{R}^{p \times p}$ and $\mathbf{X}\mathbf{Y} \in \mathbb{R}^{p \times 1}$ for each $\mathbf{X}_i, \mathbf{Y}_i$ in the dataset, and then adding it to the current Σ and Γ respectively. Because of this, at any given time, we only store $\Sigma, \Gamma, \mathbf{X}_i\mathbf{X}_i^T, \mathbf{X}_i\mathbf{Y}_i$, all of which require either $O(p^2)$ or $O(p)$ storage.

Once we have Σ and Γ , regardless of the size of B , we can calculate $\hat{\beta}_n^{\Omega^{(b)}} \in \mathbb{R}^{p \times 1}$, and then take a rolling sum of $\hat{\beta}_n^{\Omega^{(b)}}$, before taking the average and returning $\hat{\beta}_n^{ave}$. Storing each $\hat{\beta}_n^{\Omega^{(b)}}$ has storage $O(p)$, so the algorithm in total requires storage of $O(p^2)$, as desired. We do need to use storage $O(k^2)$ in order to store $\Omega \mathbb{P}_n \mathbf{X}\mathbf{X}^T \Omega^T$ in order to take its inverse, but since $k < p$, the storage of the algorithm is still of order p^2 .

Question 3

For our implementation of this function, we consider P to be a uniform distribution across the integers $\{k_{min}, k_{min} + 1, \dots, k_{max}\}$ by default, but allow the user to pass in some other callable sampling function that will sample over the same range. Likewise, Q defaults to a $N(0, 1)$ distribution or a rademacher distribution, but Q can be any callable function that takes in two arguments $(k^{(b)}, p)$ and gives a $k^{(b)} \times p$ matrix of randomly sampled values. For this implementation, this is most easily done with a `lambda` function on a NumPy function for a rademacher distribution, with the `lambda` function arguments for the size.

```

[ ]: def Randy(X, y, B = 1000, Q = lambda x,y: np.array((nr.rand(1,x*y)<.5)*2 - 1).
      ↪ reshape(x,y), k_min = 1, k_max = 100, P = None):

    n = len(y)
    p = X.shape[1]

    Sigma = np.zeros(shape = (p, p))
    Gamma = np.zeros(shape = (p, 1))
    Beta = np.zeros(shape = (p, 1))

    #First, we need to calculate Sigma and Gamma
    for i in range(0, n):
        data_temp = np.asarray(X.iloc[i,]).reshape(-1,1)
        targ_temp = y[i]
        Sigma_temp = data_temp @ data_temp.T

```

```

Gamma_temp = data_temp * targ_temp

Sigma = Sigma + Sigma_temp
Gamma = Gamma + Gamma_temp

Sigma = (1/n) * Sigma
Gamma = (1/n) * Gamma

#Now we need to calculate our B estimates for Beta
for i in range(0, B):

    #Sample k_b from P(k)
    if P is not None:
        k_b = P(k_min, k_max)
    else:
        k_b = random.randint(k_min, k_max)

    Omega = Q(k_b, p)
    Inter = Omega @ Sigma @ Omega.T
    proj_beta = Omega.T @ np.linalg.inv(Inter) @ Omega @ Gamma

    Beta = Beta + proj_beta

Beta_ave = (1/B) * Beta

return Beta_ave

```

```


[ ]: nr.seed(1)
X, y, t_coef = datasets.
    ↪make_regression(n_samples=10000,n_features=1000,n_informative=100,coef=True)
X = pd.DataFrame(X)
Beta_ave = Randy(X,y,k_min=50,k_max=100)

```

```

[ ]: plt.plot(range(1,X.shape[1]+1),Beta_ave, label = "Estimated Beta")
plt.plot(range(1,X.shape[1]+1), t_coef, label = "True Beta")
plt.legend()
plt.show()

```



Question 4

This problem is basically a projection error minimization problem. While this could be solved numerically, we can obtain a closed form solution for the optimal Ω by realizing this is just an alternative formulation of PCA.

See lectures notes from Radu Horaud (<http://perception.inrialpes.fr/~Horaud/Courses/pdf/Horaud-DAML5.pdf>) or alternatively the original chapter 12 of Bishop (2006) (<http://users.isr.ist.utl.pt/~wurmd/Livros/school/Bishop%20-%20Pattern%20Recognition%20And%20Machine%20Learning%20-%20Springer%20%202006.pdf>) for the complete derivation. Bishop (2006) showed that the solution to this minimization problem is choosing the K largest eigenvectors of the covariance X .

So our function defines function `fn` that computes the objective function, `k_dim_embedding`, which computes the eigenvectors, and then runs the simulations based on synthetic data. This true synthetic data is generated using a linear model with Gaussian noise. Then we plot a comparison between the values obtained in the objective function `fn` from both methods and then we do the same for the out-of-sample MSE.

As we can see in the first plot, by construction, the objective function takes values very close to zero for the optimal solution, while is fairly random given a random projection. Now, the interesting part is that the second plot reveal very similar out-of-sample MSEs. That is, it doesn't appear to make much of a difference using a random projection compared to the "optimal" one.

```

[ ]: import numpy as np
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
import pandas as pd

# Objective function fn
def fn(Omega, X):
    n, p = X.shape
    X_Omega = X @ Omega.T
    inner_products = X @ X.T
    Omega_inner_products = X_Omega @ X_Omega.T
    return np.sum((inner_products - Omega_inner_products)**2)

# K-dim embedding function
def k_dim_embedding(X, k):
    n, p = X.shape
    # Covariance matrix
    cov_X = np.cov(X.T)
    # Top k eigenvectors of covariance
    eigenvalues, eigenvectors = np.linalg.eigh(cov_X)
    # Take k eigenvectors corresponding to the k largest eigenvalues
    Omega_kdim = eigenvectors[:, -k:]
    # Return Omega as a (k, p) matrix
    return Omega_kdim.T

# Parameters
n = 1000 # sample size
p = 50 # number of dimensions/regressors
k = 10 # number of reduced dimensions
n_simulations = 100 # number of simulations
beta_true = np.random.normal(size=p) # true beta coefficients

# Results dataframe
results_df = pd.DataFrame(columns=['Simulation', 'Method', 'fn', 'MSE'])

for i in range(n_simulations):
    # Generate data
    X = np.random.normal(size=(n, p))
    e = np.random.normal(size=n)
    y = X @ beta_true + e

    # Split data into training and testing sets
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=0)

    # Calculate fn and MSE for random projection method
    # We choose random normal

```

```

Omega_random = np.random.normal(size=(k, p))
fn_random = fn(Omega_random, X_train)
model_random = LinearRegression()
X_train_random = X_train @ Omega_random.T
model_random.fit(X_train_random, y_train)
y_pred_random = model_random.predict(X_test @ Omega_random.T)
mse_random = np.mean((y_test - y_pred_random)**2)
results_df = pd.concat([results_df, pd.DataFrame({'Simulation': [i+1],
↪ 'Method': ['Random Projection'], 'fn': [fn_random], 'MSE': [mse_random]})])

# Calculate fn and MSE for k-dim embedding method
Omega_kdim = k_dim_embedding(X_train, k)
fn_kdim = fn(Omega_kdim, X_train)
model_kdim = LinearRegression()
X_train_kdim = X_train @ Omega_kdim.T
model_kdim.fit(X_train_kdim, y_train)
y_pred_kdim = model_kdim.predict(X_test @ Omega_kdim.T)
mse_kdim = np.mean((y_test - y_pred_kdim)**2)
results_df = pd.concat([results_df, pd.DataFrame({'Simulation': [i+1],
↪ 'Method': ['K-dim Embedding'], 'fn': [fn_kdim], 'MSE': [mse_kdim]})])

results_df.reset_index(drop=True, inplace=True)
print(results_df)

```

	Simulation	Method	fn	MSE
0	1	Random Projection	2.410688e+10	44.291845
1	1	K-dim Embedding	2.147244e+07	49.813253
2	2	Random Projection	2.166990e+10	47.610149
3	2	K-dim Embedding	2.156385e+07	44.566010
4	3	Random Projection	1.823309e+10	40.037707
..
195	98	K-dim Embedding	2.178134e+07	45.387477
196	99	Random Projection	1.728632e+10	42.036464
197	99	K-dim Embedding	2.159327e+07	49.290599
198	100	Random Projection	1.481993e+10	36.985113
199	100	K-dim Embedding	2.203076e+07	43.815145

[200 rows x 4 columns]

```

[ ]: import matplotlib.pyplot as plt

# Plot 1: fn values for each simulation
fig, ax = plt.subplots()
for key, grp in results_df.groupby('Method'):
    ax = grp.plot(ax=ax, kind='line', x='Simulation', y='fn', label=key)
ax.set_xlabel('Simulation')
ax.set_ylabel('fn')

```

```

ax.set_title('Objective function values')
plt.show()

# Plot 2: MSE values for each simulation
fig, ax = plt.subplots()
for key, grp in results_df.groupby('Method'):
    ax = grp.plot(ax=ax, kind='line', x='Simulation', y='MSE', label=key)
ax.set_xlabel('Simulation')
ax.set_ylabel('MSE')
ax.set_title('Mean squared error values (Out-of-sample)')
plt.show()

```

